

Machine Learning

Digit Classification Project

May 23, 2018

Student ID: MRPNIC004



Problem Description

In this project, we develop a classification scheme that decides whether a handwritten digit, in the form of an image, is even or odd. The idea is to use data sets consisting of pixel values (1 for black, 0 for white) for the images whereby the first column of the data set is the true value (0 to 9) of the digit shown in the image, and columns 2 to 785 are the pixel values of a given image which can be reconstructed into a 28 by 28 grid to reveal the image of the handwritten digit.

The classification schemes will be trained on a training set consisting of 2500 images which include labels of the images and tested on a set consisting of 2500 images which do not contain labels for the images. We will implement the following classification schemes:

1. Pocket Learning Algorithm
2. Logistic Regression
3. Support Vector Machine
4. Neural Network
5. Random Forests
6. Boosting
7. Regression Trees

We will compare various results of the schemes, the most important being the cross validation error. We will also compare the accuracy of the schemes for various sizes of the training and validation sets as well as compare the accuracies both with and without Principal Component Analysis applied to the training set. We will then apply the trained models to the test data set to produce predictions. We also provide various plots of the results to get more insight into the performance of the schemes.

1 Data

As mentioned above, the training set contains 2500 images all contained in a .csv format. We will denote the training set by \mathbf{X} which is a matrix of dimension 2500×785 and the output vector by \mathbf{y} which is a column vector of length 2500. Below in Figure 1, illustrates the first 6 images (rows 1-6) of the training set

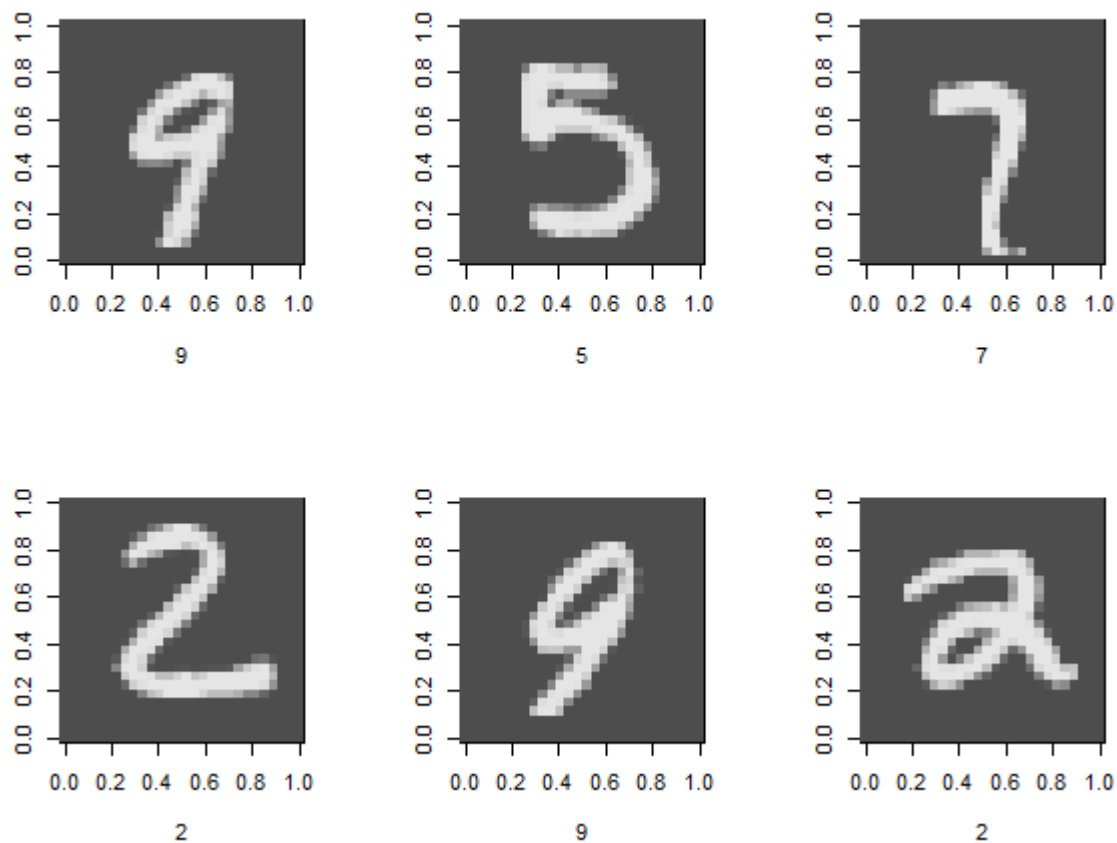


Figure 1: Illustrations of the first 6 digits in the training set with the true value of the digit shown below each image.

2 Dimension Reduction

Due to the nature of the data set, there are a large number of features/dimensions (pixels) and the data set itself consists of a large number of images. Thus, it is often useful to reduce the dimensions of the data set that we feed into our algorithm in the attempt of getting a better sense of the relationship between variables and/or features.

2.1 Principal Component Analysis

Principal Component Analysis (PCA) is one such dimension reduction method. The basic idea of PCA is to represent the input as a lower dimensional input by finding an orthonormal basis that explains most of the variance in the data set which we can use to approximate the full set of features (785) using fewer features. In this project, we will not scale the images since the pixels can only take on a value of 0 and 1 and hence scaling will not lead to drastic improvements than if we were to have images with pixels ranging between 0 and 256.

We will choose the number of principal in our orthonormal basis to be the principal components that explain 99% of our variance. This is done by finding the ratio of the cumulative sum of the eigenvalues divided sum of all eigenvalues. Once this ratio reaches 99%, we have enough principal components and in our case that number was 44.

3 Methods

3.1 Pocket Learning Algorithm

The Pocket Learning Algorithm (PLA) is a classification algorithm and is a modified version of the Perceptron Learning Algorithm. The Perceptron Learning Algorithm is designed to classify linearly separable data and hence will not terminate if the data is non-separable. The PLA is modified to 'keep in pocket' the weight which gives the lowest in sample error up to the current iteration [1]. The weight which gives the lowest in sample error throughout all of the iterations will be the final hypothesis chosen by the algorithm. The PLA is a linear model that creates a signal using a linear combination of the inputs¹. A binary function will then be applied to the signal to produce an output of 0 or 1. In this project, the PLA will iterate through the training set and pick incorrectly classified images as represented by the rows of the training set and will attempt to correctly classify the image as being odd (0) or even (1) by using the updated weights as mentioned above.

The error measure of the model will be the number of incorrectly classified images remaining after the algorithm terminates divided by the number of images in the set (2500).

3.2 Logistic Regression

The PLA above is a linear classification algorithm which produces a binary output (± 1) however if we want to get the certainty/uncertainty of this output, we can implement Logistic Regression which outputs a probability value (output $\in [0, 1]$). This is accomplished by applying a logistic function² to the signal to get an output in the range $[0, 1]$.

The error of the measure is found by showing that maximizing the likelihood of getting all the y_n 's from the corresponding x_n 's is equivalent to:

¹ $\mathbf{x}^T \mathbf{w}$ where \mathbf{x} is a vector of pixels for a given image and \mathbf{w} is the weight vector to be updated by the algorithm

²The logistic function for the signal is: $\frac{e^{\mathbf{x}^T \mathbf{w}}}{1 + e^{\mathbf{x}^T \mathbf{w}}}$

$$E_{in} = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \mathbf{w}^T \mathbf{x}})$$

The idea is to then use stochastic gradient descent to minimize this quantity.

Logistic regression performs well for linearly separable data. For non linearly separable data (as in our case) we would expect it to perform similarly to SVM.

3.3 Support Vector Machine

Support Vector Machine (SVM) is a classification scheme that attempts to find the optimal separation boundary (hyperplane) to separate the classes of the points. The key to finding the optimal boundary is maximizing the margin of the boundary. The margin is the distance between the boundary separating the points and the points that lie closest to it. The nice thing about SVM is that it is relatively simple and more importantly it copes well with non linearly separable data. It also deals well with high dimensional inputs and doesn't slow down drastically compared to other algorithms.

For non linearly separable data the idea is to allow for points which violate the boundary to be misclassified and control the number of such points with a cost parameter, C . This is known as a soft margin SVM and C is the soft-margin cost function which controls the influence of each support vector on the overall optimization. Will now have a non linear boundary which will cause problems in calculating the distance between the boundary and the closest points. This problem is solved by introducing a kernel. The kernel allows us to compute vector products of inputs from a lower dimensional space in a transformed higher dimensional space which is rendered separable. Thus, we are able to classify non linearly separable data in a robust and simple manner.

3.4 Neural Network

A neural network is made up of a series of layers where each layer consists of a series of nodes. The first layer is called the input layer where the input data is fed into the network. The hidden layers consist of weighted sums of the inputs. The weights are the parameters which we are searching for to optimize the model. The output layer then gives the output of the network which is then a weighted sum of the node values in the last hidden layer. The algorithm will feed forward the information through the network to produce an output. We then feed this error back (backpropagation) through the network to update the weights (using stochastic gradient descent) and so the process continues until we get a desired result after a number of iterations. The difficulty with neural networks is that we don't know how many nodes and layers is optimal.

3.5 Random Forests

Random Forests start by sampling points from the data set (Bootstrapped samples). We then build a decision tree from each subsample (bagged trees). However, each time a split is considered in these

trees, a random sample of mjp^3 (where p is the number of features in the training set) predictors are chosen as split candidates. This decorrelates the trees leads to better predictive performance.

3.6 Stochastic Gradient Boosting

Gradient boosting has 3 main components: a loss function to be optimized, a series of tree functions called weak learners used to make predictions and a function to sequentially add new trees to the model to reduce the loss function. Boosting tries to find the additive function $F : \mathbb{R}^d \rightarrow \mathbb{R}$ that adds the trees in such a way to minimize the loss:

$$L(F) = \sum_{i=1}^n \log(1 + \exp[-2F(x_i)\tau_i])$$

where F that minimizes $L(F)$ is given by:

$$F(x_i) = \frac{1}{2} \log \frac{\#\text{observations} : x_i, \tau = 1}{\#\text{observations} : x_i, \tau = -1} \approx \frac{1}{2} \log - \text{odds}(x)$$

Gradient descent will then be used to minimize the loss function when adding trees. Stochastic gradient descent can then be used instead of gradient descent by taking a random subsample of the training data at each iteration. The subsample is then used to fit the base learner, instead of the full data set being used [3].

4 Cross Validation

Validation (cross validation) is a key component of machine learning as it provides us with an estimate of the out-of-sample error based on information available in sample [1]. We will focus on cross-validation (R-fold cross validation) which is a modified version of validation and gives us an unbiased estimated of the out of sample error by basically taking an average of validation errors. More specifically, we partition the training data into disjoint sets, D_1, \dots, D_K , each with a size of N/K where $N = |D|$. We then train on each set D_k to get a final hypothesis g^- which is then validated on the complement of the data set. We then take an average of these K validation values to get the cross-validation error.

5 Results

5.1 Pocket Learning Algorithm

We implemented the Pocket algorithm as explained above and wrote our own function for it. The result was a cross validation error 53.4% which was really poor.

³In R, the default value is $m \approx \sqrt{p}$

5.2 Logistic Regression

We use the ‘caret’ package and use the ‘train’ function with the ‘glm’ option. Since we are doing binomial classification, we choose the bernoulli method for the function. Due to the slow computational speed of the function, we use standard validation by choosing a training set to be the first 80% of the training data and the validation set to be the last 20%. The function is slow due to the gradient descent algorithm trying to minimize the cross entropy error for a data set of high dimension.

Logistic regression had a validation error of 83.5%.

5.3 Support Vector Machine

We implemented SVM using the ‘e1071’ package in R. We tested the performance of the model by computing a 5-fold cross validation error for different combinations of cost parameters (C) and kernels. We chose 5-fold cross validation instead of 10-fold cross validation since 10-fold only improved on 5-fold by a few decimals but has much more expensive computing time. The table below shows the results of the 5-fold cross validation without PCA:

	0.01	1	5	30	80
polynomial	5.36%	5.36%	5.36%	5.36%	5.36%
linear	21.28%	21.28%	21.28%	21.28%	21.28%
radial	49.80%	49.80%	49.80%	49.80%	49.80%

Table 1: Without PCA: 5-fold cross validation errors for different values of C shown in the top of the table vs different kernels shown on the left of the table.

It is clear that the cost parameter (C) does not affect the outcome of the process for any of the kernels. It is also obvious that the polynomial kernel outperforms the linear and the radial basis kernel. We thus choose a cost parameter of 1 and the polynomial kernel for predicting on the test set.

The table below shows the results of 5-fold cross validation with PCA:

	0.01	1	5	30	80
polynomial	5.56%	5.56%	5.56%	5.56%	5.56%
linear	16.08%	42.32%	42.32%	42.32%	42.32%
radial	51.24%	51.24%	51.24%	51.24%	51.24%

Table 2: With PCA: 5-fold cross validation errors for different values of C shown in the top of the table vs different kernels shown on the left of the table.

Thus doing PCA does not improve the accuracy. From the 2 tables above, the optimal accuracy for SVM is: 94.64%

5.4 Neural Network

We implement a Neural Network using the ‘nnet’ function in R. Instead of implementing cross validation (since the computational time will be hours), we use standard validation by choosing a training set to be the first 80% of the training data and the validation set to be the last 20%. We use 1 hidden layer with 40 nodes. We chose our activation function to be the logistic function since our true output is chosen to be 0 or 1 so we can round the final output of the neural network to get predictions which we can compare to our true output.

We got a validation error 15.8% of and hence an accuracy of 84.2%. The algorithm was the slowest of all and took over 40 minutes to compute with a performance not much better than any of the other algorithms.

5.5 Random Forests

We implemented Random Forests using the ‘randomForest’ package in R. We implement 5-fold cross validation to compute the cross validation errors for different values for the number of trees used in the algorithm. The table below shows the results of Random Forest without PCA:

No. Trees	E_{cv}
10	8.88%
30	7.36%
60	6.12%
100	6.56%
200	6.48%
300	6.04%

Table 3: Without PCA: 5-fold cross validation errors for different numbers of trees used in the random forest function.

There is no drastic improvement in the cross validation error for trees larger than 60. In fact, 100 and 200 trees leads to a larger error than for 60 trees so we choose 60 trees to be our optimal number of trees to predict on the test data.

The table below shows the 5-fold cross validation errors for Random Forests with PCA:

No. Trees	E_{cv}
10	13.32%
30	10.32%
60	8.72%
100	8.48 %
200	8.32 %
300	8.28 %

Table 4: Without PCA: 5-fold cross validation errors for different numbers of trees used in the random forest function.

Thus, PCA does not improve the cross validation errors.

The optimal (optimal parameter values) accuracy for Random Forests is: 93.88%.

The figure below shows the error plot for random forest for different numbers of trees when the model is trained on the full training set. As can be seen, as we increase the number of trees past 100, there is barely an increase in the accuracy of the model.




Figure 2: Overall Out of Bag error and the errors for each of the classes (0 and 1).

5.6 Boosting

To implement Boosting, we implement the 'gbm' package in R. We implement 5-fold cross validation to choose the optimal combination of shrinkage parameter and number of trees used in the model. The table below shows the cross validation errors for the various combinations when PCA is not used:

trees/ shrinkage	0.1	0.01	0.001
100	15.08%	24.52%	27.32%
400	13.48%	18.84%	26.64%
800	13.60%	16.28%	25.36%
1500	13.80%	14.72%	22.60%

Table 5: Without PCA: 5-fold cross validation errors for different combinations of shrinkage parameters shown in the top of the table and different numbers of trees shown on the left.

It is clear that using more than 400 trees doesn't lead to smaller cross validation errors for a shrinkage parameter of 0.1 which gives us our lowest error. Hence, we choose a shrinkage parameter of 0.1 and 400 trees to train our boosting model on all of the training data to produce a prediction on the test data.

The table below shows the cross validation errors for the various combinations when PCA is used:

trees/ shrinkage	0.1	0.01	0.001
100	17.84%	26.16%	31.88%
400	14.28%	21.68%	31.80%
800	13.08%	17.80%	27.68%
1500	13.08%	15.84%	25.40%

Table 6: With PCA: 5-fold cross validation errors for different combinations of shrinkage parameters shown in the top of the table and different numbers of trees shown on the left.

The optimal (optimal parameter values) accuracy for Boosting is: 86.52%.

5.7 Recursive Partitioning Trees

We implement classification trees using the recursive partitioning tree function, 'rpart', in R. We use the 'rpart.control' function to create a stopping criterion for the tree. We choose the minimum number of observations that must exist in a node in order for a split to be attempted to be 10 and the minimum number of observations in the terminal nodes to be 5.

We get an accuracy of 84.72% without PCA and 81.76% with PCA.

Figure 3 shows the recursive partitioning tree trained on the full training set. The algorithm, as illustrated by the diagram, starts with the complete data set in the first node right at the top of the

tree and chooses the best feature to split into the next 2 nodes. In each box, the number at the top represents the way that the node has voted, the numbers below that tell us the proportion of pixels are less (left) and greater than (right) the pixel value shown below the box and the last number in the box is the percentage of the total data set used in the decision at a given node.

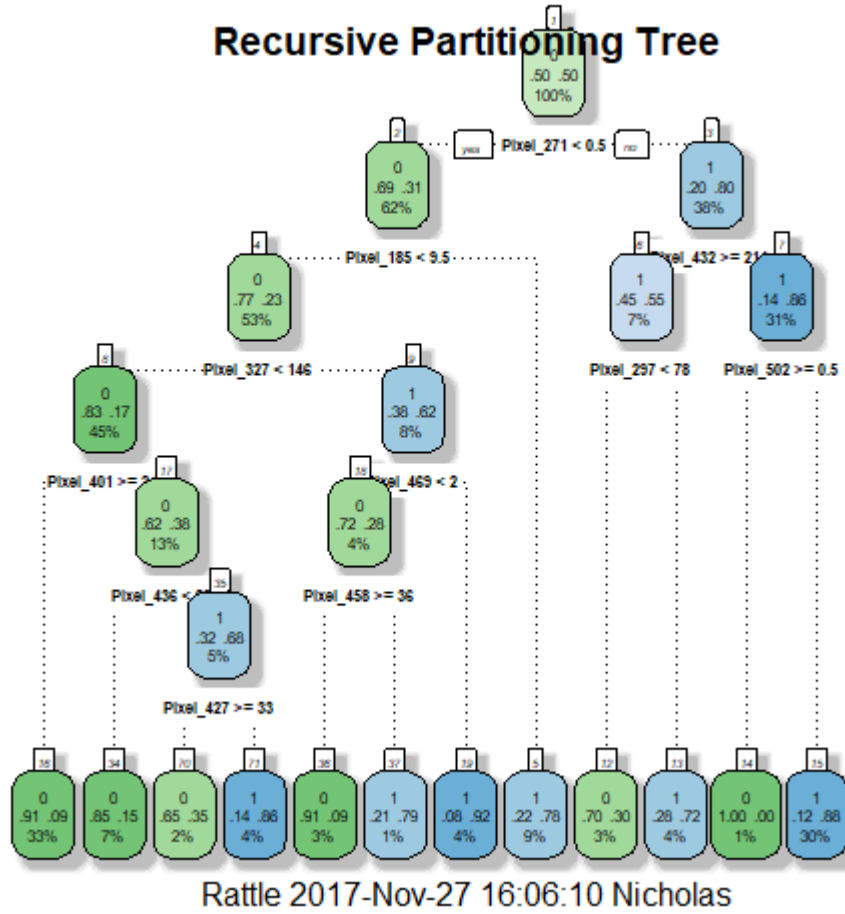


Figure 3: Recursive partitioning tree for the full training set.

6 Conclusion

It is clear that the use of PCA did not lead to drastic improvements in the accuracy of the algorithms although it did lead to slight improvements in the speed of the algorithms. SVM outperformed all of the other algorithms with an accuracy of 94.64%. The Random Forest algorithm also performs relatively well and has an accuracy of 93.88%. The neural network was the slowest algorithm of all

and took over 40 minutes to run without cross validation. Logistic Regression was also very slow due to the gradient boosting step needed to minimize the in sample (cross entropy) error. The simpler models, Boosting and Recursive Partitioning trees didnt perform as well, and had accuracies of 86.52% and 84.72% respectively.

The attached .csv file, Predictions.csv, includes the predictions of all of the studied algorithms on the test set which does not include the true values of the digits. It would be interesting to test whether SVM outperforms the other on the test set as well to conclude that it has great in sample and out of sample performance for the chose parameters.

7 References

References

- [1] Yaser S Abu-Mostafa, Malik Magdon-Ismail, & Hsuan-Tien Lin. Learning from data, volume 4. AML Book New York, NY, USA:, 2012.
- [2] Jerome Friedman, Trevor Hastie, and Robert Tibshirani. The Elements of Statistical Learning, volume 1. Springer Series in Statistics New York, 2001.
- [3] J. Friedman (1999). Stochastic Gradient Boosting.

Appendix

R Code

```
## Machine Learning Project
#
#                               Author: N. Murphy
#
## Version: Masters-Coursework-ML-ML-Project-NMURPHY.R
#
## Version Control:
# ~/R/ML/Assignments/Project/
#
# Problem Specification: This script implements Q1 of assignment 2
# Data Specification: None
# Configuration Control: userpath/MATLAB/Master/Coursework/ML
# Version Control: No version control
# References: None

## 1. Data Description
#
# None

## 2. Clear workspace
rm(list=ls()) # clear environment

## 3. Paths
# 3.1. setwd("<location of your dataset>")
rootp <- getwd()
setwd("..") # move up one level in the directory tree
# setwd(path.expand('~'))
filentrain <- "/Train_Digits_20171108.csv"
filentest <- "/Test_Digits_20171108.csv"
fpath <- "~/R/ML/Assignments/Project"
ffilentrain <- paste(fpath,filentrain,sep="")
ffilentest <- paste(fpath,filentest,sep="")

## 4. Load Train Data
traindata <- read.csv(ffilentrain)
testdata <- read.csv(ffilentest)

## 5. Load libraries
library(e1071)
library(randomForest)
library(neuralnet)
library(nnet)
library(rpart)
library(rattle)
library(gbm)
library(tree)
library(glm)

#####
```

```

## Pre-processing ##
#####

# true output
ytemp <- traindata[,1]
# Even/odd
is.even <- function(x) x%%2 == 0
is.odd <- function(x) x%%2 == 1
eveninds <- which(is.even(ytemp), arr.ind = TRUE)
oddinds <- which(is.odd(ytemp), arr.ind = FALSE)

# Convert y to +-1 for odd (-1) and even (+1)
y <- matrix(NA,nrow=length(ytemp),ncol=1)
y[eveninds,1] <- 1
y[oddinds,1] <- 0#-1

## Inputs
X <- traindata[,2:785]

## Split into training and validation (80:20 split)
# N <- dim(X)[1]
# Xtrain <- X[1:(N*0.8),]
# Xval <- tail(X,N*0.2)
# ytrain <- y[1:(N*0.8)]
# yval <- tail(y,N*0.2)
N <- dim(X)[1]
# Xtrain <- X
ytrain <- y

## Reduced training set using Principal Component Analysis
Xtrain_cov <- var(X)
PC_Xtrain <- prcomp(Xtrain_cov)

## find number of principal components to use- use number of principal components that explain at least
eigs <- PC_Xtrain$sdev^2
fv <- rbind(Cumulative = cumsum(eigs)/sum(eigs))
numPC <- which(fv[1,]>0.99)[1]

## create reduced training set
Xtrain_PCA <- as.matrix(X)%*%PC_Xtrain$rotation[,1:numPC] #extract first numpc princ. comps
Xtrain <- Xtrain_PCA

#####
## View some images from training set ##
# Create a 28*28 matrix with pixel colour values
m = matrix(unlist(traindata[10,-1]), nrow = 28, byrow = TRUE)
# Plot that matrix
image(m,col=grey.colors(255))
# reverses (rotates the matrix)
rotate <- function(x) t(apply(x, 2, rev))
# Plot the first 6 images
par(mfrow=c(2,3))

```

```

lapply(1:6,
  function(x) image(
    rotate(matrix(unlist(traindata[x,-1]),nrow = 28, byrow = FALSE)),
    col = grey.colors(256),
    xlab = traindata[x,1]
  )
)

## Implement various classification methods

#####
# 1. Pocket Algorithm ##
#####
# Vectorized version to take in NxP matrix X
iter <- 1200
E_in_w_hat <- matrix(NA,iter,1)
E_in_w <- matrix(NA,iter,1)

PLA = function(X,y,w0)
{
  X = cbind(1,X)
  w_hat = matrix(w0,dim(X)[2],1)
  misclass = (sign(as.matrix(X)%*%w_hat)!=y)
  for (i in 1:iter)
  {
    pick = sample(which(misclass==TRUE),1)
    w = w_hat +X[pick,]*y[pick]
    #evaluate errors for updates w and w_hat which has given lower E_in so far
    misclass_w = (sign(as.matrix(X)%*%w)!=y)
    misclass_w_hat = (sign(as.matrix(X)%*%w_hat)!=y)
    E_in_w[i,1] <- sum(misclass_w)
    E_in_w_hat[i,1] <- sum(misclass_w_hat)
    if (E_in_w_hat[i,1]>E_in_w[i,1] ){
      w_hat <- w
      misclass <- misclass_w
    } else {
      w_hat <- w_hat
      misclass <- misclass_w_hat
    }
  }

  # err<- (sign(as.matrix(X)%*%t(w))-y) ^2
}
return(list(E_in_w_hat,E_in_w,w_hat,w))
}

## 5-fold Cross-validation
R <- 5
folds <- cut(seq(1,nrow(Xtrain)),breaks=10,labels=FALSE)
Ein_pocket <- cbind(rep(NA,R))
E_val_pock <- Ein_pocket
# ValData <- list()

```

```

# TrainData <- list()

## Split data into 10 folds
for(j in 1:R){
  #Segement your data by fold using the which() function
  Inds <- which(folds==j,arr.ind=TRUE) # indices of current validation fold
  ValData <- Xtrain[Inds,]
  TrainData <- Xtrain[-Inds,]
  ydatatrain <- ytrain[-Inds,]
  ydataval <- ytrain[Inds,]

  res = PLA(TrainData,ydatatrain,matrix(0,dim(TrainData)[2]+1,1))

  #Ein_percept[j] <- tail(res[[2]]/2500,1)
  Ein_pocket[j] <- tail(res[[1]]/dim(TrainData)[1],1)

  ## Validation error for pocket and perceptron
  E_val_pock[j] <- sum(sign(as.matrix(cbind(1,ValData))%*%res[[3]])!=ydataval)/250
  #E_val_perc[j] <- sum(sign(as.matrix(cbind(1,ValData))%*%res[[4]])!=ydataval)/500
}

## Cross-val. error
E_CV_pock <- mean(E_val_pock)
##vector of in sample errors
Ein_pocket

## Plot insample error vs iter
plot(1:iter,res[[2]]/2500,type = "l") #Ein percept
plot(1:iter,res[[1]]/2500,type = "l") #Ein pocket

#####
# 2. Logistic Regression ##
#####

## 5-fold Cross-validation
R <- 5
folds <- cut(seq(1,nrow(Xtrain)),breaks=10,labels=FALSE)
Ein_pocket <- cbind(rep(NA,R))
E_val_pock <- Ein_pocket
# ValData <- list()
# TrainData <- list()

## Split data into 10 folds
for(j in 1:R){
  #Segement your data by fold using the which() function
  Inds <- which(folds==j,arr.ind=TRUE) # indices of current validation fold
  ValData <- Xtrain[Inds,]
  TrainData <- Xtrain[-Inds,]
  ydatatrain <- ytrain[-Inds,]
  ydataval <- ytrain[Inds,]

```



```

# Implement the glm model for logistic regression
GLM <- glm(as.formula(paste('ydatatrain ~ ',paste(paste('Pixel_',1:784,sep=''), collapse='+'), sep='
GLM <- glm(factor(ydatatrain)~.,data=TrainData ,family="binomial")

## Predict on validation data
pred_glm <- predict.glm(GLM,newdata = ValData, type="response")
#
# ## Confusion matrix
# table(`Actual Class` = yval, `Predicted Class` = pred_glm)

# Error and accuracy
E_val_glm <- sum(ydataval != pred_glm)/nrow(ydataval)
Acc_SVM = 1 - E_val_glm

}

# Fit a neural network 1 hidden layers
Xtrainglm <- X[1:(N*0.8),]
ytrainglm <- y[1:(N*0.8)]
Xval <- tail(X,N*0.2)
yval <- tail(ytrain,N*0.2)
GLM <- train(Xtrainglm ,factor(ytrainglm ),method = "glm" ,family = "binomial")
pred_glm <- predict(GLM,newdata = Xval)
E_val_glm <- sum(yval!= pred_glm)/length(yval)

## Test data prediciton
pred_glmtest <- predict(GLM,newdata = testdata[,2:785])
pred_glmtest <- ifelse(pred_glmtest,1,0)

#####
# 3. Support Vector Machine ##
#####

### Test different kernals using CV:
#different kernals
kern = c('polynomial','linear','radial')
#C = 50
## R-fold Cross-validation
R <- 5
folds <- cut(seq(1,nrow(Xtrain)),breaks=R,labels=FALSE)
E_val_svm <- cbind(rep(NA,R))
E_CV_svmkern <- cbind(rep(NA,length(kern)))

### CV error for different cost parameters
C <- c(0.01,1,5,30,80)
## R-fold Cross-validation
R <- 5
folds <- cut(seq(1,nrow(Xtrain)),breaks=R,labels=FALSE)
E_val_svm <- cbind(rep(NA,R))
E_CV_svm <- matrix(NA,length(kern),length(C))

```

```

for (k in 1:length(kern)){
  for (i in 1:length(C)){
    ## Split data into R folds
    for(j in 1:R){
      #Segement your data by fold using the which() function
      Inds <- which(folds==j,arr.ind=TRUE) # indices of current validation fold
      ValData <- Xtrain[Inds,]
      TrainData <- Xtrain[-Inds,]
      ydatatrain <- ytrain[-Inds,]
      ydataval <- ytrain[Inds,]

      ## Using R's in-built function 'svm'
      SVM_fn <- svm(factor(ydatatrain) ~ ., data=TrainData, type='C-classification', kernel=kern[k], cost=C[i])

      ## Predict on validation data
      predSVM <- predict(SVM_fn,newdata = ValData, type = "class")

      ## Confusion matrix
      table(`Actual Class` = ydataval, `Predicted Class` = predSVM)

      # Error and accuracy
      E_val_svm[j] <- sum(ydataval != predSVM)/length(ydataval)
      #Acc_SVM = 1 - E_SVM
    }

    ## Cross-validation error SVM
    E_CV_svm[k,i] <- mean(E_val_svm)
  }
}

## find cost and kernal which gave lowest CV error
indssvm <- which(E_CV_svm==min(E_CV_svm))
# number of trees
print(kern[indssvm[1]])
#cost
print(C[indssvm[1]])

## data frame of results for svm CV
df_SVM <- data.frame(100*E_CV_svm)
rownames(df_SVM) <- kern
colnames(df_SVM) <- C
library(knitr)
kable(df_SVM)

## Predict on test data
SVM_fntest <- svm(factor(ytrain) ~ ., data=X, type='C-classification', kernel=kern[indssvm[1]], cost=C[indssvm[2]])

pred_svmtest <- as.matrix(predict(SVM_fntest,newdata = testdata[,2:785]))

#####
# 4. Neural Network ##
#####

```

```

# Fit a neural network 1 hidden layers
XtrainNN <- X[1:(N*0.8),]
ytrainNN <- y[1:(N*0.8)]
Xval <- tail(X,N*0.2)
yval <- tail(y,N*0.2)
nnet_model <- nnet(XtrainNN ,ytrainNN, size = 40, MaxNWts=1000000 )
predresults <- predict(nnet_model,newdata=Xval)
results <- table(round(predresults),yval)
# ### gives sum(diag(results))/length(y) = 15.8% error

pred_NNtest <- predict(nnet_model,newdata=testdata[,2:785])
pred_NNtest <- round(pred_NNtest)
# #sum(round(Predicted)!=yval)
#results <- data.frame(actual = yval, predprob = predNN$net.result, prediction = round(predNN$net.resu

#
# ptm <- proc.time()
# NN <- neuralnet(as.formula(paste('ytrainNN ~ ',paste(paste('Pixel_',1:784,sep=''), collapse='+'), sep=
# proc.time() - ptm
#
# ## Predictions on validation set
# predNN <- compute(NN,Xval)
#
# ## Confusion matrix
# CM_NN <- table(round(predNN$net.result),yval)
#
# ## Accuracy/Validation error
# (sum(diag(CM_NN)))/sum(CM_NN)
#
# ## predictions on test data
# NN <- neuralnet(as.formula(paste('ytrain ~ ',paste(paste('Pixel_',1:784,sep=''), collapse='+'), sep=
# pred_NNtest <- compute(NN,testdata[,2:785])

#####
# 5. Random Forests ##
#####
## Cv to choose number of trees for random forest
numtreerf <- c(10,30,60,100,200,300) #number of tress to use for random forest

## R-fold Cross-validation
R <- 5
folds <- cut(seq(1,nrow(Xtrain)),breaks=R,labels=FALSE)
E_val_rf <- cbind(rep(NA,R))
E_in_rf <- cbind(rep(NA,R))
E_CV_rf <- cbind(rep(NA,length(numtreerf)))

for (i in 1:length(numtreerf)){
## Split data into 10 folds
for(j in 1:R){
#Segement your data by fold using the which() function

```

```

Inds <- which(folds==j,arr.ind=TRUE) # indices of current validation fold
ValData <- Xtrain[Inds,]
TrainData <- Xtrain[-Inds,]
ydatatrain <- ytrain[-Inds,]
ydataval <- ytrain[Inds,]

## Implement Random Forest using randomForest R package
rf <- randomForest(factor(ydatatrain) ~ .,data=TrainData,type="class",ntree=numtreerf[i],importance=T)

## In sample error
E_in_rf <- tail(rf$err.rate,1)[1]

## Predict on the validation data
rfPredval = predict(rf, newdata=ValData)

# Confusion matrix- where predicted values agree and disagree with target y
CM_rf = table(rfPredval, ydataval)

## accuracy and error of the prediction on val. data
accuracy_rf = (sum(diag(CM_rf)))/sum(CM_rf)
E_val_rf[j] = 100*(1-accuracy_rf) #validation error
}
## Cross-validation error RF
E_CV_rf[i] <- mean(E_val_rf)
}
nooftreerf <- which(E_CV_rf==min(E_CV_rf))
print(numtreerf[nooftreerf])

## data frame of results for svm CV
df_rf <- data.frame(E_CV_rf)
rownames(df_rf) <- numtreerf
colnames(df_rf) <- c("Trees","Cross Validation Error")
library(knitr)
kable(df_rf)

## Train on all data using chosen number of trees
rffinal <- randomForest(factor(ytrain) ~ .,data=X,type="class",ntree=numtreerf[nooftreerf],importance=T)

## Predict on test data
Pred_rftest <- as.matrix(predict(rffinal, newdata=testdata[,2:785]))

## plot the mean square error of the forest object as function of no. of trees
plot(1:numtreerf[nooftreerf],rffinal$err.rate[,1],type="l",col="black",xlab = "No. of Trees",ylab="Error")
lines(1:numtreerf[nooftreerf],rffinal$err.rate[,2],type="l",lty="dashed",col="green") #for classification
lines(1:numtreerf[nooftreerf],rffinal$err.rate[,3],type="l",lty="dashed",col="red") #for classification
legend("top", cex =0.5,legend=c(0,"OOB",1), colnames("Even","Odd"),lty=c(2,1,2), col=c("green","black",

#####
# 6. Generalized Boosted Regression Models ##
#####

```

```

## Implement R GBM function
# fitControl <- trainControl(method = "repeatedcv", number = 10, repeats = 1)
# GBM = train(Xtrain,ytrain, method= 'gbm', trControl=fitControl)
#
## Predict on the validation data
# pred_gbm = predict(GBM, newdata=Xval)
#
## Confusion matrix
# CM_GBM = table(round(pred_gbm), yval)
#
## accuracy and error of the prediction on val. data
# accuracy_GBM = (sum(diag(CM_GBM)))/sum(CM_GBM)
# E_val_gbm = 100*(1-accuracy_GBM) #validation error

## Predict on test data
# Pred_gbm_test <- predict(GBM, newdata=testdata)

## Use CV to find number of trees
ptm <- proc.time()
numtrees_gbm <- c(100,400,800,1500)
shrinkparam = c(0.1,0.01,0.001)
## R-fold Cross-validation
R <- 5
folds <- cut(seq(1,nrow(Xtrain)),breaks=R,labels=FALSE)
E_val_gbm <- cbind(rep(NA,R))
E_CV_gbm <- matrix(NA,length(numtrees_gbm),length(shrinkparam))

for (k in 1:length(shrinkparam)){
  for (i in 1:length(numtrees_gbm)){
    ## Split data into 10 folds
    for(j in 1:R){
      #Segement your data by fold using the which() function
      Inds <- which(folds==j,arr.ind=TRUE) # indices of current validation fold
      ValData <- data.frame(Xtrain[Inds,])
      TrainData <- data.frame(Xtrain[-Inds,])
      ydatatrain <- ytrain[-Inds,]
      ydataval <- ytrain[Inds,]

      boost <- gbm(ydatatrain ~ ., data=TrainData, distribution="bernoulli", n.trees=numtrees_gbm[i],shrinkage=shrinkparam[k])
      proc.time() - ptm

      pred_GBM = predict(boost, newdata=ValData,n.trees = numtrees_gbm[i],type = "response")

      CM_GBM = table(round(pred_GBM),ydataval)

      ## accuracy and error of the prediction on val. data
      accuracy_GBM = (sum(diag(CM_GBM)))/length(ydataval)
      E_val_gbm[j] = 100*(1-accuracy_GBM) #validation error
    }
    ## Cross-validation error RF
    E_CV_gbm[i,k] <- mean(E_val_gbm)
  }
}

```

```

}
proc.time() - ptm
## find lowest cv error for given number of trees
indsgbm <- which(E_CV_gbm==min(E_CV_gbm))
# number of trees
print(numtrees_gbm[indsgbm[2]])
#shrinkage
print(shrinkparam[indsgbm[1]])

## data frame of results for svm CV
df_gbm <- data.frame(E_CV_gbm)
rownames(df_gbm) <- numtrees_gbm
colnames(df_gbm) <- shrinkparam
library(knitr)
kable(df_gbm)

## Predict on test set using full training data with chosen number of trees from CV

boost <- gbm(as.formula(paste('ytrain ~ ',paste(paste('Pixel_',1:784,sep=''), collapse='+'), sep='')),
pred_GBMtest = as.matrix(round(predict(boost, newdata=testdata[,2:785],n.trees =numtrees_gbm[indsgbm[2]]

#####
# 7. Classification trees ##
#####

## Stopping criteria
stopcrit <- rpart.control(minbucket = 5, minsplit = 10)

## cross validation
R <- 5
E_val_tree <- cbind(rep(NA,R))
E_CV_tree <- cbind(rep(NA,R))
## Split data into R folds
for(j in 1:R){
  #Segement your data by fold using the which() function
  Inds <- which(folds==j,arr.ind=TRUE) # indices of current validation fold
  ValData <- data.frame(Xtrain[Inds,])
  TrainData <- data.frame(Xtrain[-Inds,])
  ydatatrain <- ytrain[-Inds,]
  ydataval <- ytrain[Inds,]

  ## Implement 'rpart' function
  fulltree <- rpart(ydatatrain ~ ., method = "class", data = TrainData,control = stopcrit)

  ## Predict on the validation data
  pred_tree <- predict(fulltree, newdata = ValData, type = "class")
  E_val_tree[j] <- sum(ydataval != pred_tree)/length(ydataval)
}
E_CV_tree <- 100*mean(E_val_tree)

## Implement 'rpart' function on full data set

```

```

stopcrit <- rpart.control(minbucket = 5, minsplit = 10)
fulltree <- rpart(ytrain ~ ., method = "class", data = X, control = stopcrit)
## Plot the tree
# plot(fulltree)
# text(fulltree, cex = 0.5)
library(rattle)
fancyRpartPlot(fulltree, tweak=1.5, main="Recursive Partitioning Tree")

## Predict on test data
pred_treetest <- as.matrix(predict(fulltree, newdata = testdata[,2:785], type = "class"))

#####
### Predictions on test data ##
Predictions <- cbind(pred_svmtest, Pred_rfctest, pred_GBMtest, pred_treetest, pred_NNtest, as.matrix(pred_glmtest))
Pred <- data.frame(Predictions)
colnames(Pred) <- c("SVM", "Random Forest", "Boosting", "Tree", "Neural Network", "Logistic Regression")

## Write to .csv
write.csv(Pred, "~/R/ML/Assignments/Project/Predictions.csv")

#write.csv(pred_svmtest, "~/R/ML/Assignments/Project/Predictions.csv")

### EOF

```